

1. A. 00 0010 0111 1100

B.

| | |
|-------------|------|
| VPN | 0x9 |
| TLBI | 0x1 |
| TLBT | 0x2 |
| TLB hit? | N |
| page fault? | N |
| PPN | 0x17 |

C. 0101 1111 1100

D.

| | |
|-------------|------|
| CO | 0x0 |
| CI | 0xf |
| CT | 0x17 |
| cache hit? | N |
| cache byte? | - |

2. A. 00 0011 1010 1001

B.

| | |
|-------------|------|
| VPN | 0xe |
| TLBI | 0x2 |
| TLBT | 0x3 |
| TLB hit? | N |
| page fault? | N |
| PPN | 0x11 |

C. 0100 0110 1001

D.

| | |
|-------------|------|
| CO | 0x1 |
| CI | 0xa |
| CT | 0x11 |
| cache hit? | N |
| cache byte? | - |

3. A. 00 0000 0100 0000

B.

| | |
|-------------|-----|
| VPN | 0x1 |
| TLBI | 0x1 |
| TLBT | 0x0 |
| TLB hit? | N |
| page fault? | Y |
| PPN | - |

C. n/a

D. n/a

-
4.
 - (a) First we must read the PTE, which is at index $VA.VPN$ (0x08063) or at address $PTBR + 0x08063 * 4$, or $PTE[PTBR + 0x2018C]$.
 - (b) Next we check fields P (must be 1) and S (must be 0) of that PTE (no additional memory reads). If either of these are not as required, we are done.
 - (c) Finally we compute the physical address, which is $(PTE[PTBR + 0x2018C].PFN \ll 12) + VA.PO$ ($= 0x472$), and read from that physical address.
 5.
 - (a) First we must read the PDE, which is at index $VA[31:22]$ (0x020) or at address $PTBR + 0x020 * 4$, or $PTE[PTBR + 0x080]$.
 - (b) Next we check fields P (must be 1) and S (must be 0) of that PDE (no additional memory reads). If either of these are not as required, we are done.
 - (c) Next we must read the PTE, which is at index $VA[21:12]$ (0x063) or at address $PDE[PTBR + 0x080].PFN \ll 12 + 0x063 * 4$, or $PTE[PDE[PTBR + 0x080].PFN \ll 12 + 0x18C]$.
 - (d) Next we check fields P (must be 1) and S (must be 0) of that PTE (no additional memory reads). If either of these are not as required, we are done.
 - (e) Finally we compute the physical address, which is $(PTE[PDE[PTBR + 0x080].PFN \ll 12 + 0x18C].PFN \ll 12) + VA.PO$ ($= 0x472$), and read from that physical address.
 6. Both set associative caches and virtual memory want to replace the item that will be used longest in the future, but this can't be implemented. The best approximation is LRU (Least Recently Used), which virtual memory can use (or approximate with multiple FIFOs) because the VM system has lots of time and wants to work hard to avoid a miss. A cache must use a cheaper approximation because it can't afford the overhead and so normally uses random.