

CPSC 313, Summer 2013 — Final Exam **Solution**

Date: June 24, 2013; Instructor: Mike Feeley

1 (10 marks) This question is concerned with the general principles of pipelining, dependencies and hazards.

First, consider two hardware implementations, **S** and **D**, of the same ISA. They have different pipelines: **S**'s pipeline has 5 stages and **D**'s has 7. Answer the following two questions. You must use the term *instruction-level parallelism* in one of your answers below.

1a Give one reason why **S** might execute some programs faster than **D**. Explain carefully.

Higher CPI.

1b Give one reason why **D** might execute some programs faster than **S**. Explain carefully.

Faster clock.

Now, consider the dependencies that exist in the following code snippets. List **every** dependency. For each of them: (a) list the type of dependency and the register or memory location involved (if appropriate); (b) say whether the dependency is a hazard in the Y86 pipeline; and (c) if it is a hazard, explain how the hazard is resolved in the standard Y86-Pipe implementation (with data forwarding and predicting jumps are taken) and how many stalls are required; if it is not a hazard, explain why not.

1c

```
mrmovl    (%eax), %ebx
addl      %ebx, %ebx
```

Causal on %ebx is (load-use) hazard that is resolved using data forwarding with one stall.
Output %ebx is not a hazard because size state writes registers and so they can not be re-ordered.

1d

```
rmmovl    %eax, (%ebx)
mrmovl    (%ebx), %eax
```

Anti on %eax is not a hazard because register read comes earlier in pipeline than register write and so they can not be re-ordered.
Causal on memory location whose address is stored in %ebx is not a hazard because there is only one stage that reads/writes memory and so they can not be re-ordered.

1e

```
jle      foo
```

Control hazard resolved using branch prediction results in no stalls if jump is taken and two stalls if it is not.

2 (10 marks) In the standard Y86-Pipe implementation, the load-use hazard requires one stall to resolve. But there are some forms of this hazard that should not actually require this stall. Consider, for example, this code snippet that performs a memory-to-memory copy, copying a value from the stack into memory at the address in %ebx.

```
popl      %eax
rmmovl    %eax, (%ebx)
```

2a In the standard pipeline, where there is a stall, the new value of register %eax is forwarded from the first instruction to the second. List the stage it is forwarded from and say whether it is forwarded from the beginning or end of that stage. List the stage it is forwarded to and say whether it is forwarded to the beginning or end of that stage.

From end of M to end of D.

2b Explain in plain English why the pipeline could be modified to eliminate the stall in this case, but not in the case of certain other load-use hazards.

The value of %eax is not needed by the rmmovl instruction until it reaches the memory stage where as other instructions that read registers need the value one stage earlier, in the execute stage.

2c Fixing the pipeline to eliminate this unnecessary stall requires changing the way that register %eax is forwarded. List the stage it is forwarded from in the revised pipeline and say whether it is forwarded from the beginning or end of that stage. List the stage it is forwarded to and say whether it is forwarded to the beginning or end of that stage.

From end of M to end of E.

- 2d** Now, show the *changes* to the forwarding logic needed to implement this change. Use the simulator's Java syntax. As you did in the simulator, prefix pipeline-stage input registers with the capital letter of stage that reads them and output registers with the lower-case letter of the stage that writes them. Use `get()` and `getValueProduced()` to read pipeline registers and `set()` to write them.

```
if (E.iCd==I_RMMOVL && M.iCd==I_POPL && E.srcA == M.dstM)
    e.valA = M.valM
```

- 2e** What other instructions could have their load-use hazards resolved in exactly this way? List instructions that could be in the first location (the load) and then instructions that could be in the second location (the use).

load instructions:

```
mrmovl (rA) and ret (%esp)
```

use instructions:

```
pushl (rA) and rrmovl (rA)
```

- 3 (20 marks)** In this question you will describe the implementation of a new Y86 instruction that conditionally jumps to an address that is stored in memory: `jxx D(rA)`. This instruction is similar to `jxx Dest` except that the jump-target address is read from memory (`M[D + r[rA]]`) instead of being a constant stored in the instruction. Use `I_JXX_DI` if you need to refer to the new instruction's op-code (i.e., `iCd`).

For example the following code jumps to address `0x1000`.

```
irmovl $0x1000, %eax    # eax = jump target address
rmmovl %eax, 4(%ebx)    # store jump target address in memory at address 4+r[ebx]
xorl    %ecx, %ecx      # ALU operation that stores zero in %ecx
je      4(%ebx)          # goto m[4+r[ebx]] if last ALU resulted in 0, which it did
```

As a starting point, consider the following implementation of the standard `jxx Dest` instruction for Y86-PipeMinus that you will recall handles hazards by stalling and does not do data forwarding or jump prediction. (*This is repeated on a back page so that you can remove it.*)

FETCH.select_pc:

```
if (M.iCd.get() == I_JXX && M.cnd.get() == 0)
    f.pc.set (M.valP.get());
else
    f.pc.set (F.prPc.get());
```

FETCH:

```
f.iCd.set (mem.read (f.pc.getValueProduced()+0, 1) [0] .value() >>> 4);
f.iFn.set (mem.read (f.pc.getValueProduced()+0, 1) [0] .value() & 0xf);
f.valC.set (mem.readIntegerUnaligned (f.pc.getValueProduced()+2));
f.valP.set (f.pc.getValueProduced()+5);
```

FETCH.predict_pc:

```
f.prPc.set (f.valC.getValueProduced());
```

DECODE:

```
d.valA.set (R_NONE);
d.valB.set (R_NONE);
d.dstE.set (R_NONE);
d.dstM.set (R_NONE);
```

EXECUTE:

```
e.cnd.set (read-condition-codes-and-opcode-to-decide-if-jump-is-taken);
```

PIPELINE HAZARD CONTROL:

```
if (D.iCd.get() == I_JXX || E.iCd.get() == I_JXX ) {
    F.stall = true;
    D.bubble = true;
}
```

Answer these questions about the implementation of the new instruction, `jxx D(rA)`.

- 3a** Provide the following parts of the Y86-PipeMinus implementation of the new instruction (just the new instruction; i.e., ignore all of the other instructions in the ISA just as was done in the example above). Use the same syntax constraints as in Question 2d.

FETCH_select_pc:

```
if (W.iCd.get() == I_JXX_DI && W.cnd.get() == 1)
    f.pc.set (W.valM.get());
else
    f.pc.set (F.prPc.get());
```

FETCH_predict_pc:

```
f.prPc.set (f.valP.getValueProduced());
```

EXECUTE:

```
e.valE.set (E.valC.get() + E.valA.get());
e.cnd.set (read-condition-codes-and-opcode-to-decide-if-jump-is-taken);
```

MEMORY:

```
m.valM.set (mem.readInteger (M.valE.get()));
```

PIPELINE HAZARD CONTROL:

```
if (D.iCd.get() == I_JXX_DI || E.iCd.get() == I_JXX_DI ||
    M.iCd.get() == I_JXX_DI)
{
    F.stall = true;
    D.bubble = true;
}
```

- 3b** Now consider the Y86-Pipe implementation of this instruction where your goal is to stall as little as possible by using static jump prediction (only static jump prediction and not using things like jump caches that implement dynamic jump prediction). Describe your strategy in plain English and explain how this strategy differs from the one used for the standard `jxx Dest`, if it does, and why. Be sure to say whether the mis-prediction penalty changes and if so to what and why.

Predict not taken instead of taken, because we don't know the target address in F. Checking for mis-prediction must wait an additional cycle since we don't know the target address until the instruction completes M and has read the address from memory. Thus, mis-prediction penalty is 3 bubbles.

- 3c** Draw a diagram of the pipeline state as the new instruction moves through the pipeline from D to W in the case where the pipeline mis-predicts the jump and then corrects this mis-prediction. Your diagram should have 4 lines.

- 3d** Provide the following parts of the Y86-Pipe implementation of the new instruction. Be careful to prevent a mis-predicted instruction from changing system state (i.e., registers, memory or condition codes).

FETCH_select_pc:

```
if (W.iCd.get() == I_JXX_DI && W.cnd.get() == 1)
    f.pc.set (W.valM.get());
else
    f.pc.set (F.prPc.get());
```

FETCH_predict_pc:

```
f.prPc.set (f.valP.getValueProduced());
```

PIPELINE HAZARD CONTROL:

```
if (E.iCd.get() == I_JXX_DI && e.cnd.getValueProduced() == 1) {
    F.stall = true;
    D.bubble = true;
    E.bubble = true;
}
if (M.iCd.get() == I_JXX_DI && M.cnd.getValueProduced() == 1) {
    F.stall = true;
    D.bubble = true;
}
```

4 (12 marks) The following questions relate to the performance of programs running on the standard Y86 Pipe processor with data forwarding and predicting that branches are taken.

4a What two things must you measure to compute a program's throughput on a pipelined processor?

Clock rate and CPI.

4b Consider the following program running on the standard Y86-Pipe processor (with data forwarding and predicting that conditional jumps are taken).

```

    irmovl $1, %eax      # eax = 1
    irmovl $4, %ebx      # ebx = 4
    irmovl $1000, %ecx   # ecx = 1000
    irmovl 0x2000, %edi   # edi = 0x2000
    irmovl $0, %esi      # esi = 0
L0: subl    %eax, %ecx    # ecx = ecx - 1
    je      L1           # goto L1 if ecx = 0
    addl    %ebx, %edi    # edi = edi + 4
    mrmovl  (%edi), %ebp  # ebp = m [edi]
    addl    %ebp, %esi    # esi = esi + ebp
    jmp     L0           # goto L0
L1: rmmovl  %esi, (%edi)  # m [edi] = esi

```

What is the CPI (average cycles per instruction) of the loop (i.e., between L0 and the jmp L0, inclusive). Justify your answer and show your work by annotating the code to indicate where stalls occur.

The loop contains 6 instructions. The body of the loop is executed 1000 times; the first two instructions are executed 1001 times. And so, the total number of instruction executions is 6002. There is one load-use hazard that adds one bubble to each iteration. The conditional jump is mis-predicted 1000 times and correctly prediction once; each mis-prediction adds 2 bubbles. And so, the total number of bubbles added is $1000 + 2000 = 3000$ and so the total number of cycles is $3000 + 6002 = 9002$. And thus $CPI = \frac{9002}{6002}$, which is approximately 1.5.

4c Re-write the code starting with L0 to improve its pipeline performance.

```

    addl    %ebx, %edi
    subl    %eax, %ecx
    je      L1
L0: mrmovl  (%edi), %ebp
    addl    %ebx, %edi
    addl    %ebp, %esi
    subl    %eax, %ecx
    jne     L0
L1: rmmovl  %esi, (%edi)

```

4d What is the CPI of your revised code

The revised loop has 5 instructions, four of which are executed 1000 times; the last instruction is executed 1001 times. And so, the total number of instruction executions is 5001. There are no load-use hazards that require stalling and the condition jump is predicted correctly 1000 times and incorrectly once, adding two bubbles. And so the total number of cycles is 5003 and thus $CPI = \frac{5003}{5001}$, which is approximately 1.

5 (10 marks) Answer the following questions about the configuration of a 8-MB (i.e., 2^{23} -B), 8-way set associative cache with 64-B (i.e., 2^6 -B) blocks, using 32-bit memory addresses. Recall that only data blocks are included when computing the size of the cache. Assume that address bits are numbered 0 to 31, with 0 the least-significant bit. For questions c-e, give an a bit range (e.g., "bits 3 to 5").

5a How many blocks does this cache have?

$$\frac{8 \times 2^{20}}{64} = \frac{2^{23}}{2^6} = 2^{17}$$

5b How many sets does this cache have?

$$\frac{2^{17}}{8} = \frac{2^{17}}{2^3} = 2^{14}$$

5c Which address bits store the tag?

Bits 20–31.

5d Which address bits store the index?

Bits 6–19.

5e Which address bits store the block offset?

Bits 0–5.

6 (10 marks) Fill in the following table describing a program’s behaviour in a direct-mapped cache where: **bits 0-3 of the address store the block offset** and **bits 4-7 store the index**. The program performs the memory accesses listed in the table below, and only these, in the order listed. The array `a`, which starts at address `0x1000` is declared as `“int a[16][16];”`.

Access	Hex Address	Tag	Index	Hit?	Explain Miss
<code>a[0][0]</code>	<code>0x1000</code>	<code>0x10</code>	0	N	compulsory
<code>a[2][3]</code>	<code>0x108c</code>	<code>0x10</code>	8	N	compulsory
<code>a[0][2]</code>	<code>0x1008</code>	<code>0x10</code>	0	Y	
<code>a[4][1]</code>	<code>0x1104</code>	<code>0x11</code>	0	N	compulsory
<code>a[0][3]</code>	<code>0x100c</code>	<code>0x10</code>	0	N	conflict
<code>a[4][1]</code>	<code>0x1104</code>	<code>0x11</code>	0	N	conflict

7 (10 marks) In this question you consider the tradeoffs involved in changing a single feature of a cache design at a time. Note that only the listed feature changes (e.g., the overall size of the cache does not change). Carefully answer each question. **Justify your answers.**

7a What is the main benefit of increasing block size?

Increased cache hits for programs with spatial locality.

7b What is the main benefit of decreasing block size?

Increased cache hits with programs with temporal locality, but limited spatial locality.

7c What is the main benefit of increasing associativity?

Decreased number of conflict misses.

7d What is the main benefit of a write-back cache compared to a write-through cache with write buffer?

Write absorption.

7e What is the main benefit of a write-through cache with write buffer compared to a write-back cache?

No need for write-allocate penalty on write miss.

8 (3 marks) Disk and file system performance.

8a Give a formula for disk-read latency in terms of the following parameters

- N = number of bytes read
- S = maximum seek time (in units of seconds)
- R = rotation speed of disk (in units of rotations per second)
- D = track density (in units of bytes per track)

Assume average seek and rotational delay of $\frac{1}{2}$ of maximum.

$$\frac{S}{2} + \frac{R}{2} + \frac{N}{D \times R}$$

9 (9 marks) In a system that uses the Clock Algorithm for virtual-memory page replacement, explain the purpose of each of the following *page-table-entry (PTE)* bits by describing how they are cleared, set to 1 and read. Say whether the operation is performed by hardware or software, what causes it and why it is done. Each answer can be a single sentence of the form (where “/” indicates alternatives): “Hardware/Software when xxx to indicate/determine that/whether yyy.”. For example, if the question asked what set’s the PFN, a good answer would be: ”Software when a page is faulted from disk into memory to record the physical location of the page for subsequent virtual-to-physical address translations.”

9a valid

- set to 0:

Software when page is invalidated due to replacement or when process terminates.

- set to 1:

Software when page is added to page table.

- read:

Hardware to determine whether page-table entry is valid; successful translation if 1; page fault if 0.

9b accessed

- set to 0:

Software when moved to inactive list to indicate that page has not been accessed.

- set to 1:

Hardware when page is accessed.

- read:

Software when looking for a replacement victim at the bottom of inactive list; page is replaced if its 0 and moved to top of active list if its 1.

9c dirty

- set to 0:

Software when dirty page is written back to disk.

- set to 1:

Hardware when page is modified.

- read:

Software to determine whether page is dirty, to write page to disk or when looking for a replacement victim; dirty pages can not be replaced.

10 (6 marks) File and VM replacement.

10a Briefly explain what LRU replacement is and how it relates to the theoretical optimal replacement algorithm.

Replace the page that was last accessed furthest in the past. Optimal is to replace the page accessed furthest in the future. The last access time of a page is used as a heuristic to predict next access time; i.e., pages accessed long ago are less likely to be accessed in near future than pages accessed recently.

10b Give an example of a file-system workload (i.e., a program that reads data a file) where LRU replacement would be a poor choice.

Sequential access to a large file.

10c Explain how application-level buffering is different from caching and how it can be effective in managing memory used to store disk data in some cases where caching and LRU replacement fail.

Caching relies on locality to improve performance; keeping recently accessed blocks in the cache. Buffering is a explicit technique that involves reading ahead from a file (for example) so that data is in the buffer when needed. Buffering is particularly useful (and easy) for sequential access. Using buffering, only the next few, yet-to-be-read pages are kept in the buffer. Once a page has been read, its buffer location can be immediately re-used. This approach is much better than the cache for this workload, since the cache would end up keeping the whole file in the cache, which, since there is no temporal locality, is pointless.

